

# Recurrent Neural Networks

Benjamin Roth

Centrum für Informations- und Sprachverarbeitung  
Ludwig-Maximilian-Universität München  
`beroth@cis.uni-muenchen.de`

# Recursive Neural Networks: Motivation

How do you ...


- ... best represent a sequence of words as a vector?
- ... combine the learned word vectors effectively?
- ... retain the information relevant to a particular task (certain features of particular words), suppress unessential aspects?



# Recursive Neural Networks: Idea


- Calculate for each position ( “ time step ” ) in the text a representation that summarizes all essential information up to this position.
- For a position  $t$  this representation is a vector  $\mathbf{h}^{(t)}$  (hidden representation)
- $\mathbf{h}^{(t)}$  is calculated recursively from the word vector  $\mathbf{x}^{(t)}$  and the hidden vector of the previous position:

$$\mathbf{h}^{(t)} = f(\mathbf{h}^{(t-1)}, \mathbf{x}^{(t)})$$



# Recursive Neural Networks

$$\mathbf{h}^{(t)} = f(\mathbf{h}^{(t-1)}, \mathbf{x}^{(t)})$$



- The hidden vector in the last time step  $\mathbf{h}^{(n)}$  can then be used for classification ( “ *Sentiment of the sentence?* ” )
- The predecessor representation of the first time step uses the  $\mathbf{0}$  vector (containing only zeros).

# Recursive function $f$

$$\mathbf{h}^{(t)} = f(\mathbf{h}^{(t-1)}, \mathbf{x}^{(t)})$$

- The  $f$  function takes two vectors as input and outputs a vector.
- The function  $f$  is in most cases a combination of:
  - ▶ **Vector matrix multiplication:**
  - ▶ and a **non-linear function** (e.g., logistic sigmoid) applied to all components of the resulting vector.

$$\mathbf{h}^{(t)} = \sigma(\mathbf{W}[\mathbf{h}^{(t-1)}; \mathbf{x}^{(t)}] + \mathbf{b})$$

Usually a bias vector  $\mathbf{b}$  is added, which is sometimes omitted for simplicity.

# Recursive function $f$

$$\mathbf{h}^{(t)} = f(\mathbf{h}^{(t-1)}, \mathbf{x}^{(t)})$$

## • Vector matrix multiplication:

- ▶ Simplest form of mapping a vector onto a vector.
- ▶ First, the vectors  $\mathbf{h}^{(t-1)}$  ( $k$  components) and  $\mathbf{x}^{(t)}$  ( $m$  components) are concatenated:
  - ★ Result  $[\mathbf{h}^{(t-1)}; \mathbf{x}^{(t)}]$  has  $k + m$  components.
- ▶ Weight matrix  $W$  (size:  $k \times (k + m)$ )
  - ★ the same matrix for all time steps (*weight sharing*)
  - ★ is optimized when training the RNN.

# Recursive function $f$

$$\mathbf{h}^{(t)} = f(\mathbf{h}^{(t-1)}, \mathbf{x}^{(t)}) = \sigma(\mathbf{W}[\mathbf{h}^{(t-1)}; \mathbf{x}^{(t)}] + \mathbf{b})$$

- **Non-linear function**


- ▶ Examples: Sigmoid, Tanh (= scaled sigmoid, between  $-1 \dots 1$ ), Softmax, ReLu ( $= \max(0, x)$ )
- ▶ Applied to all components of the resulting vector.
- ▶ Necessary so that the network can compute interesting, non-linear interactions, such as the effect of negation.



# Neural Networks: Terminology

# Layers


- Conceptually, a neural network is composed of several (*layers*).
- Each layer is a function that takes a vector (or matrix) as the input, and outputs a vector (or matrix).
- The size of the output does not have to match the size of the input (also vector  $\leftrightarrow$  matrix possible).
- The output of the previous layer is the input for the next layer.



**Which layers are there in our example (prediction of sentiment with RNN)?**

# Layers predicting sentiment with (simple) RNN

- Input: vector with word-ids
- Layer 1 (Embedding): Lookup of word vectors for ids (vector→matrix)
- Layer 2 (RNN): Calculation of the sentence vector from word vectors (matrix→vector)
- Layer 3: Calculation of the probability for positive sentiment from the sentence vector (vector→Real number, represented as a vector with 1 element)

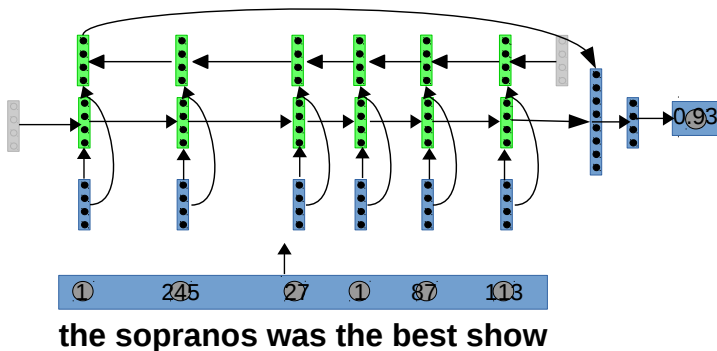


**the sopranos was the best show**

Outlined in red: inputs / outputs


# Prediction with RNN: Possible extensions (1)

- A second RNN can process the sentence from right to left: The two RNN representations are then concatenated.



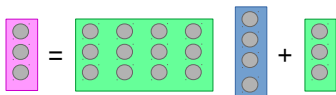
## Prediction with RNN: Possible extensions (2)

- Before the prediction, several *Dense* layers can be cascaded.
  - ▶ A dense layer (also: *fully connected layer*) corresponds to a matrix multiplication (+ bias) and application of a non-linearity
  - ▶ A Dense layer “translates” vectors and combines information from the previous layer.
  - ▶ Usually, the prediction layer is a dense layer. (in the example: translation into a vector of size 1, nonlinearity is the sigmoid function)




# Dense-Layer: illustration

- $\mathbf{y} = \sigma(\mathbf{W}\mathbf{x} + \mathbf{b})$ 
  - ▶  $\mathbf{W}$  and  $\mathbf{b}$  are parameters that have to be learned by the model
  - ▶ The nonlinearity  $\sigma$  is applied element by element
- $\hat{\mathbf{y}} = \mathbf{W}\mathbf{x} + \mathbf{b}$



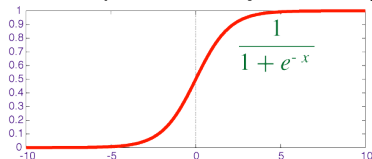
- $\mathbf{y} = \sigma(\hat{\mathbf{y}})$



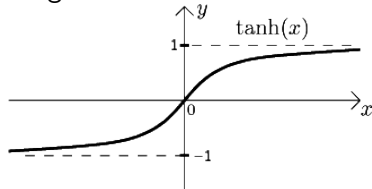
**Note: In a simple RNN, the recursive function corresponds to a dense layer!**

# Frequently used nonlinearities

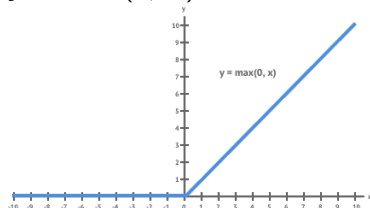
- Logistic Sigmoid:  $y_i = \sigma(\mathbf{x}_i)$   
Value range between 0 ... 1, can be interpreted as a **probability**.



- Tanh:  
 $y_i = \tanh(\mathbf{x}_i) = 2\sigma(2\mathbf{x}_i) - 1$   
Like Logistic Sigmoid, but value range between -1 ... 1



- ReLU (*rectified linear unit*):  
 $y_i = \max(0, \mathbf{x}_i)$



- Softmax:

$$y_i = \frac{e^{(x_i)}}{\sum_j e^{(x_j)}}$$

- Normalizes the output of the preceding layers to a **probability distribution**
- Mostly used in output layer for prediction

## Note on learning the model parameters

- A neural network is a function built from simple units, with one vector as the input (e.g., word ids of a sentence), and another vector as the output (e.g., probability for positive sentiment).
- For a data set, a cost function can now be calculated, e.g. the negative log likelihood:
  - ▶ (negative log) probability that the model assigns to the annotated labels of the data set.
  - ▶ Sometimes also called **cross-entropy**.
- The parameters can then be optimized (similar to Word2Vec) with Stochastic Gradient Descent.
  - ▶ Parameters are e.g. Word Embeddings, Weight Layers of Dense Layers, ... etc.
  - ▶ Unlike Word2Vec, NN usually performs a parameter update on a *mini-batch* of 10-500 training instances.
  - ▶ Several extensions of SGD are available (RMS-Prop, Adagrad, Adam, ...)



# Neural Networks: Implementation with Keras

# Introduction

## What is Keras?

- Neural Network library written in Python
- Designed to be minimalistic & straight forward yet extensive
- Built on top of TensorFlow

## Keras strong points:

- Easy to get started, powerful enough to build serious models
- Takes a lot of work away from you.
- Reasonable defaults (e.g. weight matrix initialization).
- Little redundancy. Architectural details are inferred when possible (e.g. input dimensions of intermediate layers, masking).
- highly modular; easy to expand

# Keras: Idea

```
from keras.models import Sequential
from keras.layers import SomeLayer, OtherLayer
model = Sequential()
model.add(SomeLayer(...))
model.add(OtherLayer(...))
model.add(...)
model.compile(optimizer='sgd',
              loss='binary_crossentropy',
              metrics=['accuracy'])
model.fit(x_train, y_train)
```

- `Sequential()` creates a model in which layers can be sequentially stacked on each other.
  - ▶ For each layer, the corresponding object is first created and added to the model.
  - ▶ The added layer take over the output of the previous layer as its input.

# Keras: Idea

```
from keras.models import Sequential
from keras.layers import SomeLayer, OtherLayer
model = Sequential()
model.add(SomeLayer(...))
model.add(OtherLayer(...))
model.add(...)
model.compile(optimizer='sgd',
              loss='binary_crossentropy',
              metrics=['accuracy'])
model.fit(x_train, y_train)
```

- `model.compile`: When the specification of the model is completed, it can be compiled:
  - ▶ It is specified which learning algorithm should be used.
  - ▶ Which cost function should be minimized.
  - ▶ And what additional metrics should be calculated for evaluation.
- `model.fit`: Training (adjust the parameters in all layers)

# Keras: Embedding Layer

```
from keras.layers import Embedding
...
model.add(Embedding(input_dim=10000, output_dim=50))
...
```

- Provides word vectors of size `output_dim` for a vocabulary of size `input_dim`.
  - ▶ Often the first layer in a model.
  - ▶ Input per instance: vector with word id's
  - ▶ Output per instance: matrix; sequence of word vectors.
- The parameters (word vectors) of the embedding layer
  - ▶ ... can be initialized with pre-trained vectors (Word2Vec), or at random.
  - ▶ ... if you use pre-trained word vectors, further optimization of the word vectors is sometimes not necessary.

```
from keras.layers import Embedding
...
model.add(Embedding(input_dim=10000, output_dim=50, \
                    weights=[word_vectors], trainable=False))
...
```

- Advantages / disadvantages of using pre-trained word vectors and not optimizing them further?

- Advantages / disadvantages of using pre-trained word vectors and not optimizing them further?
- *Advantage: For a specific task, such as Sentiment analysis, often comparatively little training data is available. Word vectors can be trained unsupervised on large corpora, these therefore have a **better coverage**. In addition, the model has fewer parameters to optimize, which is why **there is less risk of overfitting**.*
- *Disadvantage: The word vectors used may not fit the task, the relevant properties were not taken into account in the unsupervised learning of the vectors ⇒ **Underfitting***
- *Note: A good middle ground is often to initialize the vectors with pre-trained vectors, and still further optimize them on the task-specific training data.*

## Keras: RNN Layer

- Although the previously introduced variant of the RNN is an expressive model, the parameters are difficult to optimize (*vanishing gradient problem*).
- Extensions of the RNN, which facilitate the optimization of the parameters, are e.g. **LSTM** (long short-term memory network) and **GRU** (gated recurrent unit network)

```
from keras.layers import LSTM, Bidirectional
```

```
...
```

```
model.add(LSTM(units=100))
```

```
...
```

- Two RNNs (left-to-right and right-to-left). output are the concatenated end vectors (as in the example above):
- Instead of the end vector, a matrix can also be output which contains the state vector  $h$  for each position:

```
model.add(LSTM(units=100, return_sequences=True))
```

**For which computer linguistic tasks is it necessary to have access to the state vector at each position?**



# Keras: RNN Layer

- Instead of the end vector, a matrix can also be output which contains the state vector  $h$  for each position: **For which computer linguistic tasks is it necessary to have access to the state vector at each position?**

*Whenever a prediction needs to be made for each position, e.g. part of speech tagging.*

# Keras: Dense Layer

Two options:

- As an intermediate layer
  - ▶ Combines information from previous layers.
  - ▶ Nonlinearity is ReLu or Tanh.

```
from keras.layers import Dense
...
model.add(Dense(100, activation='tanh'))
...
```

- As output layer
  - ▶ Probability of an output.
  - ▶ Non-linearity is sigmoid (probability of output 1-vs-0) or softmax (any number of classes, one-hot-encoding).

```
...
model.add(Dense(1, activation='sigmoid'))
...
```

# Training

```
model.compile(loss='binary_crossentropy', optimizer='adam',\n              metrics=['accuracy'])
```

- Loss functions:
  - ▶ `binary_crossentropy` if only one class is predicted (sigmoid activation)
  - ▶ `categorical_crossentropy` if probability distribution over several classes (Softmax activation)
- Optimizer: `adam`, `rmsprop`, `sgd`

# Training

`model.fit(...)`

Other arguments:

- Hyper-parameters
  - ▶ `batch_size`: how many instances should be used for one optimization step. (Optimization step  $\neq$  training iteration)
  - ▶ `epochs`: How many training iterations should be performed.
  - ▶ ...
- `validation_data`: Tuple (`features_dev`, `labels_dev`)  
Development data, e.g. to monitor training progress.

# Prediction and evaluation

- `y_predicted = model.predict(x_dev)`
- `score, acc, ... = model.evaluate(x_dev, y_dev)`  
Returns the value of the objective function and the metrics (loss or metrics of `model.compile`)

# Hints

- In order to be productive with Keras, it is important to become familiar with the API / Documentation!
- <https://keras.io/getting-started/sequential-model-guide/>
- Keras expects inputs as numpy arrays. Lists of various lengths (e.g., sentence representations) can be converted to a numpy array of a given number of columns by the `pad_sequences(list_of_lists, max_length)` command. (Too long lists are truncated, shorter ones are filled with 0 values) <sup>1</sup>

---

<sup>1</sup>Modul `keras.preprocessing.sequence`

# Convolutional Neural Networks

- CNNs can be used just as easily as RNNs.
- For example, to generate a CNN with 50 filters (output dimensions) and filter width 3 words for sentiment prediction ...
- ... instead of the line `model.add(LSTM (...))`, a CNN with max pooling must be used:

```
...  
model.add(Conv1D(filters=50, kernel_size=3, \  
                 activation='relu', padding='same'))  
model.add(GlobalMaxPooling1D())  
...
```

# Summary

- RNNs: Creates a sequence of vectors (*hidden states*).
- Each hidden vector is calculated recursively from the previous vector, and the word-embedding of the current position.
- A sequence may e.g. be represented by the last hidden vector.